

Model Compression with Adversarial Robustness: A Unified Optimization Framework

Shupeng Gui_{1,†}, Haotao Wang_{2,†}, Haichuan Yang₁, Chen Yu₁, Zhangyang Wang₂, and Ji Liu₃

¹Department of Computer Science, University of Rochester

²Department of Computer Science and Engineering, Texas A&M University

³Ytech Seattle AI lab, FeDA lab, AI platform, Kwai Inc

† *Equal contribution*

NeurIPS 2019, Vancouver CANADA



UNIVERSITY *of* ROCHESTER

Summary

- Goals
 - Model Compression
 - Adversarial Robustness
 - Accuracy
- **Adversarially Trained Model Compression (ATMC) Framework**
 - Adversarial objective
 - Integrating three compression ways simultaneously (Pruning, Factorization, and Quantization)
- Experiments: Better trade-off among model size, accuracy, and robustness, over currently available alternatives in various settings.



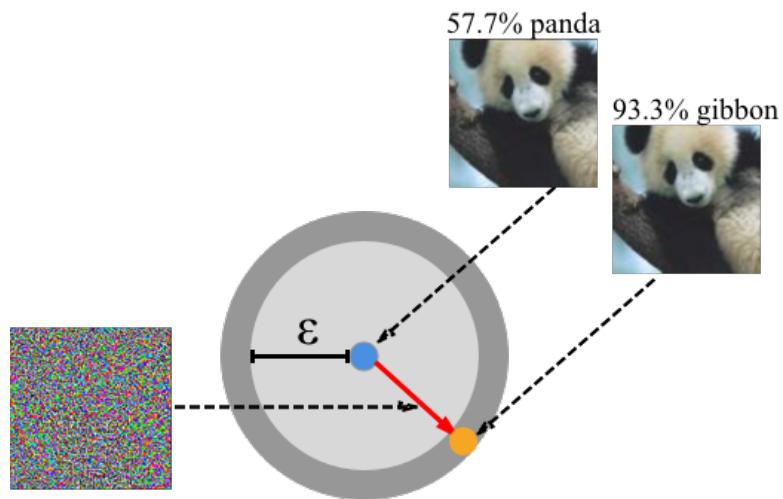
Highlights of Contributions

- First framework jointly optimizing
 - Model compression
 - Adversarial robustness
- First framework unifies all existing compression methods
 - Pruning
 - Factorization
 - Quantization

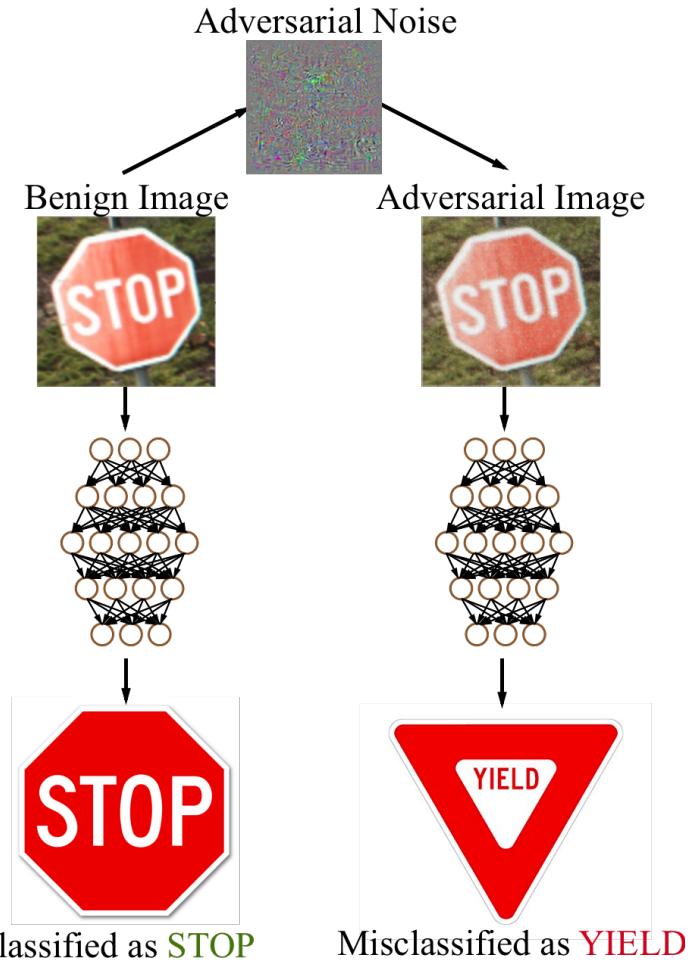


Why Adversarial Robustness?

- Easy to fool normal DNN classifiers



Goodfellow et al, "Explaining and Harnessing Adversarial Examples", ICLR 2015.

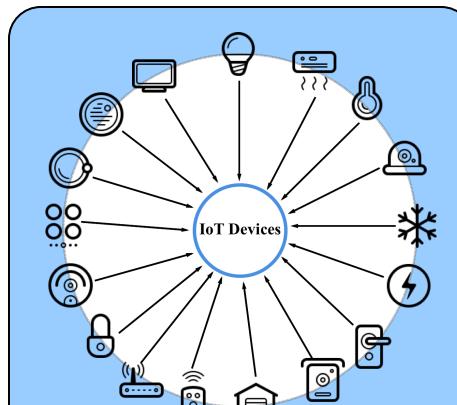


Why Model Compression?

- Efficient inference on edging devices



App & Battery



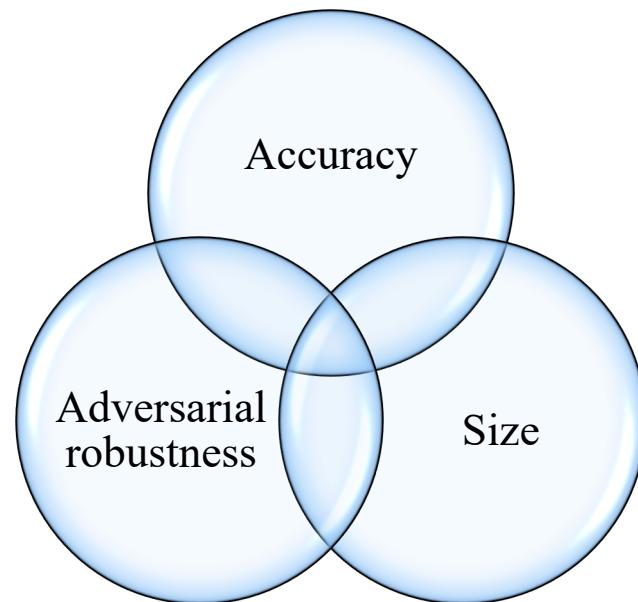
IoT & Limitation



UNIVERSITY *of* ROCHESTER

<https://www.cultofmac.com/492544/iphone-8-pack-lasers-improved-ar-autofocus/>
<https://www.digitaltrends.com/mobile/how-to-save-battery-life-on-your-smartphone/>
<https://images.app.goo.gl/92RrppKfD3bUYEwX8>

Our Work: Optimize Three Goals Simultaneously



Our Framework: Adversarially Trained Model Compression

$$\min_W \sum_{(x,y) \text{ in data set}} f^{\text{adv}}(W; x, y) \quad \text{Accuracy + Robustness}$$

$$\text{s. t. } \sum_l \|U^{(l)}\|_0 + \|V^{(l)}\|_0 + \|C^{(l)}\|_0 \leq k \quad \text{Model size}$$

$$W \in \mathcal{Q}_b := \left\{ W : \|U^{(l)}\|_0 \leq 2^b, \|V^{(l)}\|_0 \leq 2^b, \|C^{(l)}\|_0 \leq 2^b, \forall l \in [L] \right\}$$

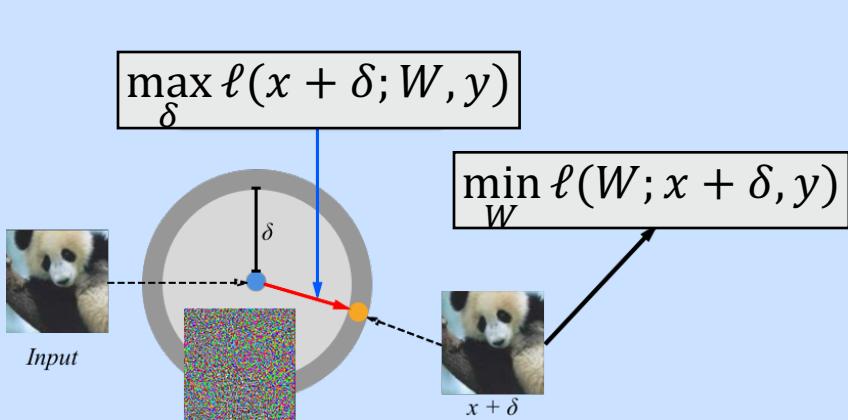
Def: $W := \{W^{(l)}\}_{l \in [L]}$, $W^{(l)} = U^{(l)}V^{(l)} + C^{(l)}$

$\|\cdot\|_0$: L-0 norm

$|\cdot|_0$: # Unique scalars



Adversarial Training Loss



$$\min_W \sum_{(x,y) \text{ in data set}} f^{\text{adv}}(W; x, y)$$

$$\text{s.t. } \sum_l \|U^{(l)}\|_0 + \|V^{(l)}\|_0 + \|C^{(l)}\|_0 \leq k$$

$$\theta \in \mathcal{Q}_b := \left\{ \theta : |U^{(l)}|_0 \leq 2^b, |V^{(l)}|_0 \leq 2^b, |C^{(l)}|_0 \leq 2^b, \forall l \in [L] \right\}$$

Def: $f^{\text{adv}}(W; x, y) = \max_{x' \in B_{\infty}^{\Delta}(x)} f(W; x', y)$

$$B_{\infty}^{\Delta}(x) := \{x' \mid \|x' - x\|_{\infty} \leq \Delta\}$$



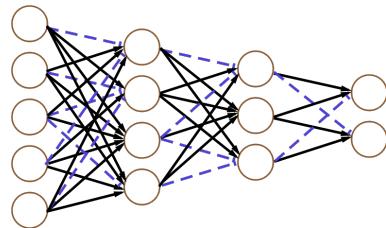
Model Size Compression: Combine Three Compression Ways

$$\min_W \sum_{(x,y) \text{ in data set}} f^{\text{adv}}(W; x, y)$$

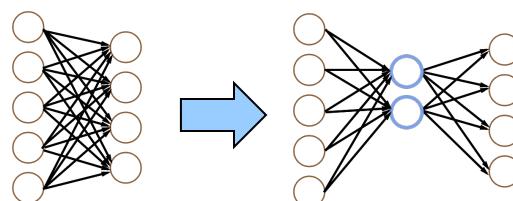
$$W := \{W^{(l)}\}_{l \in [L]}, W^{(l)} = U^{(l)}V^{(l)} + C^{(l)}$$

$$\text{s. t. } \sum_l \|U^{(l)}\|_0 + \|V^{(l)}\|_0 + \|C^{(l)}\|_0 \leq k$$

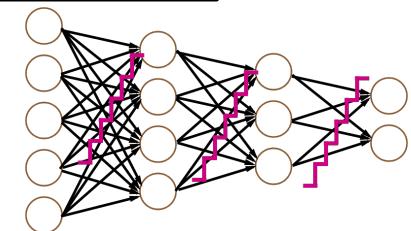
$$W \in \mathcal{Q}_b := \left\{ W : \|U^{(l)}\|_0 \leq 2^b, \|V^{(l)}\|_0 \leq 2^b, \|C^{(l)}\|_0 \leq 2^b, \forall l \in [L] \right\}$$



Weight
Pruning



Factorization



Quantization



ATMC: Optimization

- Duplicate Variables

- $\min_{\|W\|_0 \leq k, W' \in \mathcal{Q}_b} \sum_{(x,y) \text{ in data set}} f^{\text{adv}}(W; x, y) \text{ s.t. } W = W'$

- $\|W\|_0 := \sum_l \|U^{(l)}\|_0 + \|V^{(l)}\|_0 + \|C^{(l)}\|_0$

- Removing the equality constraint $W = W'$

- $\min_{\|W\|_0 \leq k, W' \in \mathcal{Q}_b} \max_u \sum_{(x,y) \text{ in data set}} f^{\text{adv}}(W; x, y) + \rho \langle u, W - W' \rangle + \frac{\rho}{2} \|W - W'\|_F^2$

- $\rho > 0$ as predefined positive number in ADMM



ATMC: Optimization

- E.g., given U in an arbitrary layer
- Update u :
 - $u_{t+1} = u_t + (U - U')$
- Update x^{adv} :
 - $x^{\text{adv}} \leftarrow \text{Proj}_{\|x' - x\|_\infty \leq \Delta} \{x + \alpha \nabla_x f(W; x, y)\}$
- Update U :
 - $U \leftarrow \text{Proj}_{\|U''\|_0 \leq k} \{U - \gamma \nabla_U [f(U; x^{\text{adv}}, y) + \frac{\rho}{2} \|U - U' + u\|_F^2]\}$
- Update U' :
 - Solving the following projection problem for U'
 - $\min_U \|U' - (U + u)\|_F^2, \quad s.t. |U'|_0 \leq 2^b$
 - Lloyd's algorithm



Experiment: CNNs

- Datasets and Benchmark Models
 - Four popular image classification datasets
 - Pick one top-performer CNN model on each

Models	#Parameters	Bit width	Model Size (bits)	Dataset & Accuracy
LeNet	430K	32	13,776,000	MNIST: 99.32%
ResNet34	21M	32	680,482,816	CIFAR-10: 93.67%
ResNet34	21M	32	681,957,376	CIFAR-100: 73.16%
WideResNet	11M	32	350,533,120	SVHN: 95.25%



Experiment: Settings

- Evaluation Metrics
 - Classification Accuracy on both benign and adversarial perturbed test set
- Model Size
 - # Non-zero elements \times the bit-width of each layer
- Compression Ratio
 - the ratio between the compressed and original model sizes
- We apply PGD attack mainly for the test of adversarial robustness



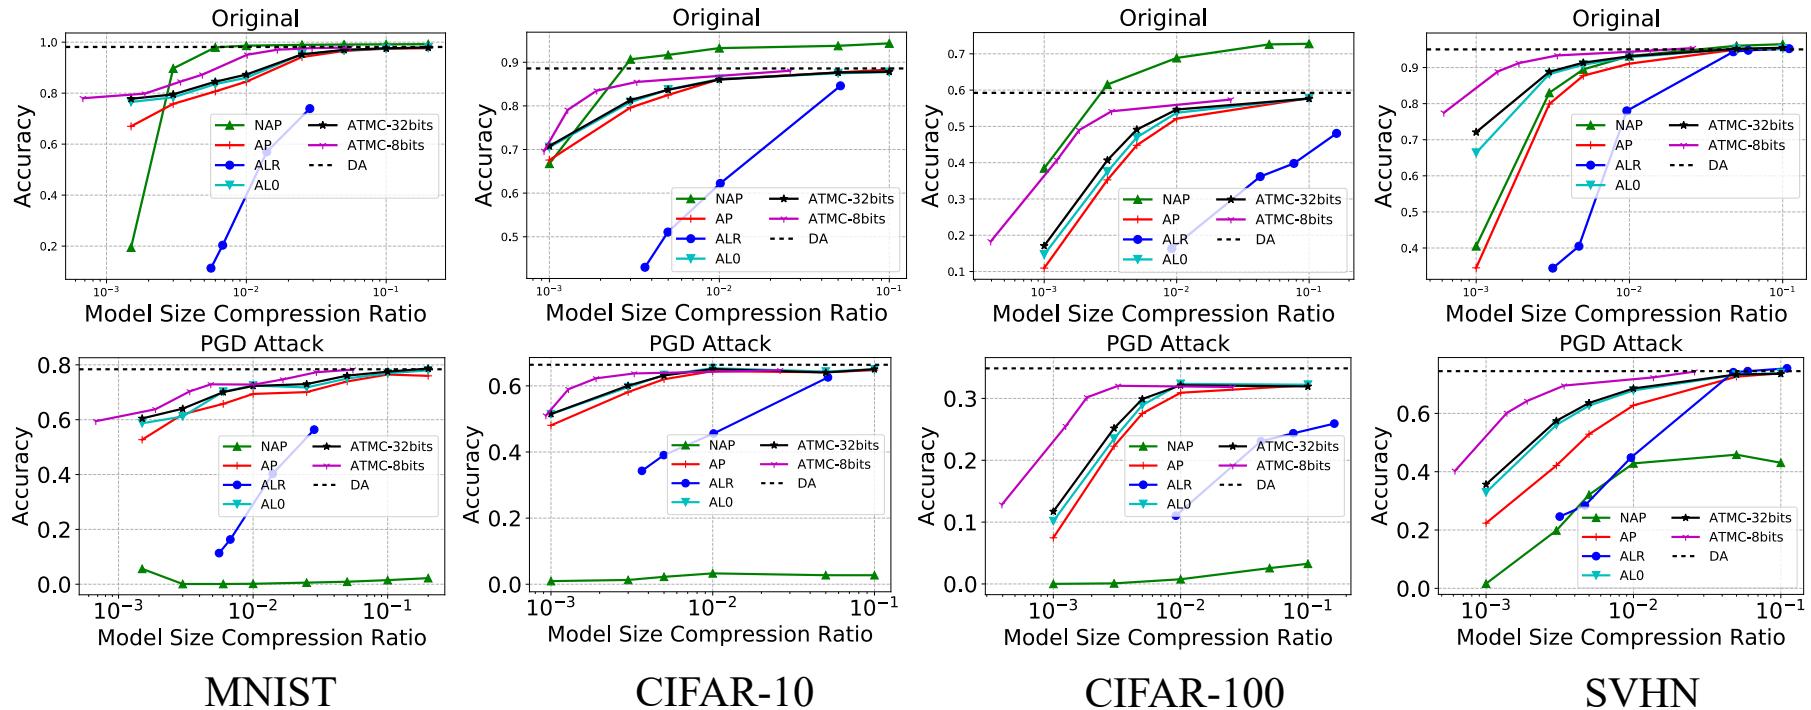
Experiment: Baselines and Alternatives

- Non-adversarial Pruning (NAP) [Han et al., NIPS 15]:
 - Pure pruning method without adversarial training
- Dense Adversarial Training (DA) [Madry et al., ICLR 18]:
 - Pure adversarial training without model compression
- Adversarial Pruning (AP):
 - NAP + Adversarial Training
- Adversarial l0 Pruning (A10)
 - L0-projection based Pruning + Adversarial Training
- Adversarial Low-Rank Decomposition (ALR)
 - Low-rank weight decomposition + Adversarial Training
- ATMC (8bits, 32 bits)
 - Our method with two bit-width settings



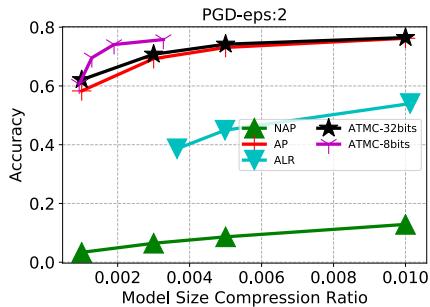
Experiment: Results

- Outstanding Performance on Trade-off between Compression and robustness for ATMC

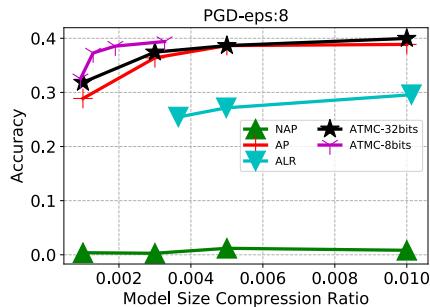


Experimental results

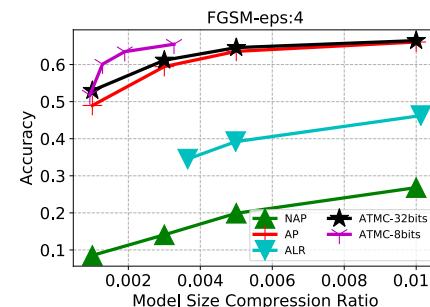
- Consistent Adversarial Robustness under various attack settings
 - Different perturbation magnitude, e.g., 2, 8
 - Different adversarial attack methods, e.g., FGSM, WRM



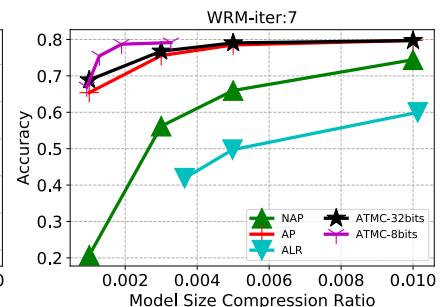
PGD,
Perturbation=2



PGD,
Perturbation=8



FGSM,
Perturbation=4



WRM,
Penalty=1.3,
Iteration=7

- More details <https://github.com/TAMU-VITA/ATMC>

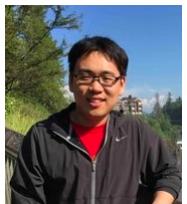


Conclusion

- ATMC, First algorithmic framework, optimizing:
 - Model Compression
 - Adversarial Robustness
- Unifying the existing compression methods:
 - Pruning
 - Factorization
 - Quantization
- Effectiveness of ATMC
 - General outstanding trade-off between model compression and robustness
 - Consistent robustness under various adversarial settings



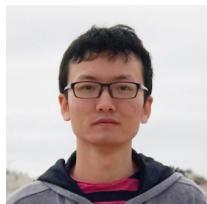
Thanks



Shupeng Gui



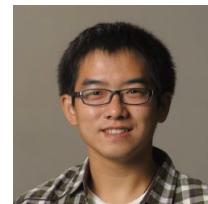
Haotao Wang



Haichuan Yang



Chen Yu



Zhangyang Wang



Ji Liu



UNIVERSITY *of* ROCHESTER

References

- Guo, Yiwen, et al. "Sparse dnns with improved adversarial robustness." *Advances in neural information processing systems*. 2018.
- Goodfellow, Ian J., Jonathon Shlens, and Christian Szegedy. "Explaining and harnessing adversarial examples." *arXiv preprint arXiv:1412.6572* (2014).
- Madry, Aleksander, et al. "Towards deep learning models resistant to adversarial attacks." *arXiv preprint arXiv:1706.06083* (2017).
- Han, Song, et al. "Learning both weights and connections for efficient neural network." *Advances in neural information processing systems*. 2015.
- Xu, Huan, Constantine Caramanis, and Shie Mannor. "Sparse algorithms are not stable: A no-free-lunch theorem." *IEEE transactions on pattern analysis and machine intelligence* 34.1 (2011): 187-193.
- Tsipras, Dimitris, et al. "Robustness may be at odds with accuracy." *arXiv preprint arXiv:1805.12152* (2018).
- Li, Hao, et al. "Pruning filters for efficient convnets." *arXiv preprint arXiv:1608.08710* (2016).
- Lloyd, "Least squares quantization in pcm", IEEE transactions on information theory, 1982
- Han, Song, Huizi Mao, and William J. Dally. "Deep compression: Compressing deep neural networks with pruning, trained quantization and huffman coding." *arXiv preprint arXiv:1510.00149* (2015).
- He, Yihui, Xiangyu Zhang, and Jian Sun. "Channel pruning for accelerating very deep neural networks." *Proceedings of the IEEE International Conference on Computer Vision*. 2017.
- Zhou, Hao, Jose M. Alvarez, and Fatih Porikli. "Less is more: Towards compact cnns." *European Conference on Computer Vision*. Springer, Cham, 2016.

